



# **Operating Systems**

## **Scheduling Algorithms**

Seyyed Ahmad Javadi

[sajavadi@aut.ac.ir](mailto:sajavadi@aut.ac.ir)

Spring 2022

---

# First-Come, First-Served



# First- Come, First-Served (FCFS) Scheduling

<u>Process</u>	<u>Burst Time</u>
$P_1$	24
$P_2$	3
$P_3$	3

- Suppose that the processes arrive in the order:  $P_1, P_2, P_3$
- The Gantt Chart for the schedule is:



- Waiting time for  $P_1 = 0$ ;  $P_2 = 24$ ;  $P_3 = 27$
- Average waiting time:  $(0 + 24 + 27)/3 = 17$

# FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order:

$P_2, P_3, P_1$

- The Gantt chart for the schedule is:



- Waiting time for  $P_1 = 6$ ;  $P_2 = 0$ ;  $P_3 = 3$
- Average waiting time:  $(6 + 0 + 3)/3 = 3$
- **Much better than previous case**

# FCFS Scheduling and Convoy effect

---

- **Short process behind long process.**
  - Consider one CPU-bound and many I/O-bound processes.



- **What is the important side-effect?**

# FCFS Scheduling and Convoy Effect (Cont.)

---

- Short process behind long process.
  - Consider one CPU-bound and many I/O-bound processes.



- What is the side-effect?
  - Results in **lower CPU and device utilization** than might be possible if the shorter **processes were allowed to go first.**

---

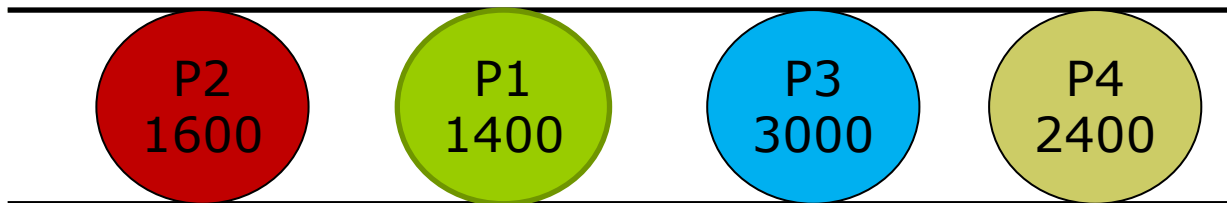
# Shortest-Job-First



# Shortest-Job-First (SJF) Scheduling

---

- **Associate** with each process the length of its **next CPU burst**.
  - Use these lengths to schedule the process with the shortest time.



Ready Queue



# Shortest-Job-First (SJF) Scheduling (cont.)

---


- **SJF is optimal**
  - Gives minimum average waiting time for a given set of processes.
- Preemptive version called **shortest-remaining-time-first**
- **The difficulty is knowing the length of the next CPU request**
  - Could ask the user
  - Estimate (we do not cover this in the class and the exams)



# Example of SJF

<u>Process</u>	<u>Burst Time</u>
$P_1$	6
$P_2$	8
$P_3$	7
$P_4$	3

## ■ SJF scheduling chart



## ■ Average waiting time = $(3 + 16 + 9 + 0) / 4 = 7$

---

# Round-Robin



# Round Robin (RR)

---

- Each process gets a small unit of CPU time (**time quantum  $q$** )
  - Usually 10-100 milliseconds.
  - After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are  $n$  processes in the ready queue and the time quantum is  $q$ :
  - Each process gets  $1/n$  of the CPU time in chunks of at most  $q$  time units at once.
  - No process waits more than  $(n-1)q$  time units.



# Round Robin (RR) (cont.)

---

- Timer *interrupts* every quantum to schedule next process
- Performance
  - $q$  large  $\Rightarrow$  FIFO
  - $q$  small  $\Rightarrow q$  must be large with respect to context switch, otherwise overhead is too high



# Example of RR with Time Quantum = 4


<u>Process</u>	<u>Burst Time</u>
----------------	-------------------

$P_1$	24
-------	----

$P_2$	3
-------	---


$P_3$	3
-------	---

- The Gantt chart is:




- Typically, higher average turnaround than SJF, but better **response**
- $q$  should be large compared to context switch time
  - $q$  usually 10 milliseconds to 100 milliseconds,
  - Context switch < 10 microseconds

# Time Quantum and Context Switch Time



# Turnaround Time Varies With The Time Quantum



process	time
$P_1$	6
$P_2$	3
$P_3$	1
$P_4$	7

A rule of thumb is that  
80% of CPU bursts should  
be shorter than  $q$



---

# Priority Scheduling



# Priority Scheduling

---

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority  
(**smallest integer  $\equiv$  highest priority**)
  - Preemptive
  - Nonpreemptive
- SJF is priority scheduling where priority is the ***inverse of predicted next CPU burst time***



# Priority Scheduling

---

## ■ Problem ≡ Starvation

- Low priority processes may never execute

## ■ Solution ≡ Aging


- As time progresses increase the priority of the process



# Example of Priority Scheduling

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
$P_1$	10	3
$P_2$	1	1
$P_3$	2	4
$P_4$	1	5
$P_5$	5	2

## ■ Priority scheduling Gantt Chart




## ■ Average waiting time = 8.2

# Priority Scheduling w/ Round-Robin

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
$P_1$	4	3
$P_2$	5	2
$P_3$	8	2
$P_4$	7	1
$P_5$	3	3

- Run the process with the highest priority. Processes with the same priority run round-robin
- Gantt Chart with time quantum = 2




---

# Multilevel Queue Scheduling



# Multilevel Queue

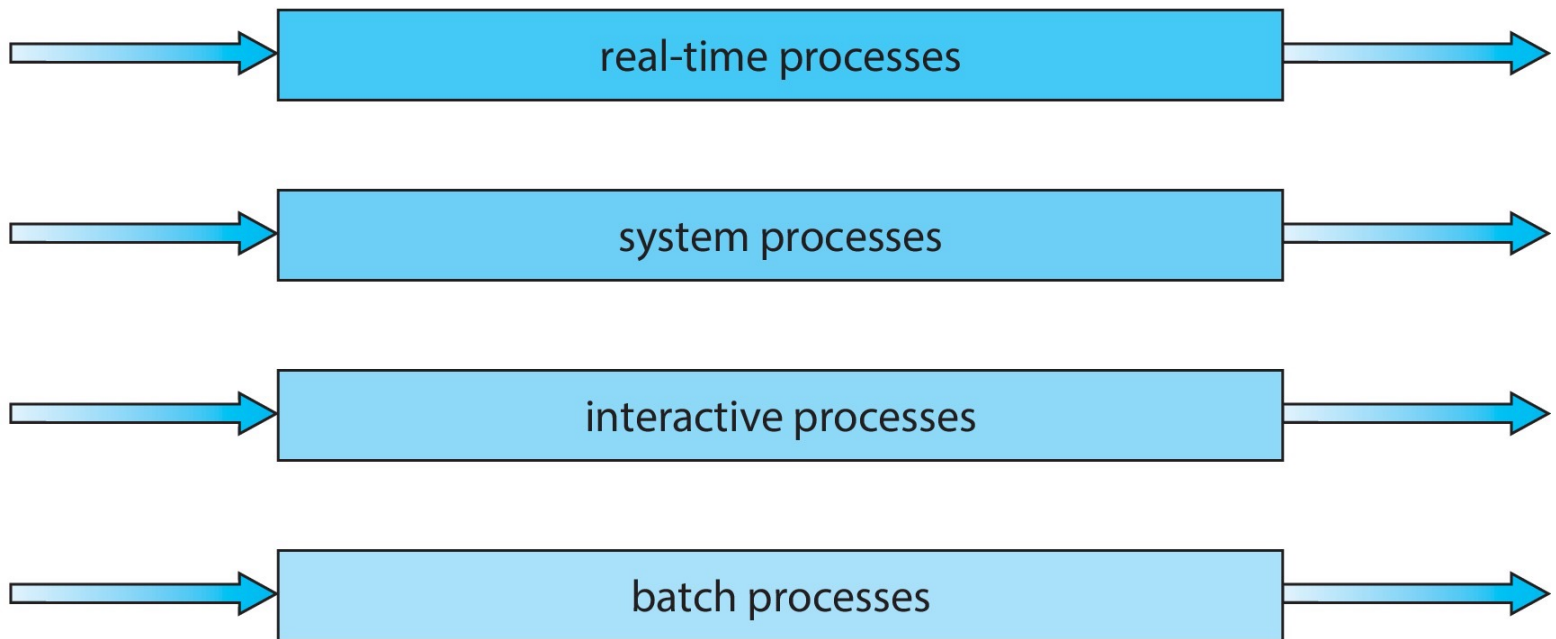
- With priority scheduling, have separate queues for each priority.
- Schedule the process in the highest-priority queue!



# Multilevel Queue

- Prioritization based upon process type

highest priority



lowest priority



# Multilevel Feedback Queue


---

- A process can move between the various queues.
- Multilevel-feedback-queue defined by the following parameters:
  - Number of queues
  - Scheduling algorithms for each queue
  - Method used to determine when to upgrade a process
  - Method used to determine when to demote a process
  - Method used to determine which queue a process will enter when that process needs service
- **Aging** can be implemented using multilevel feedback queue



# Example of Multilevel Feedback Queue

- Three queues:
  - $Q_0$  – RR with time quantum 8 milliseconds
  - $Q_1$  – RR time quantum 16 milliseconds
  - $Q_2$  – FCFS



# Example of Multilevel Feedback Queue (cont.)

## ■ Scheduling

- A new process enters queue  $Q_0$  which is served in RR
  - ▶ When it gains CPU, the process receives 8 milliseconds
  - ▶ If it does not finish in 8 milliseconds, the process is moved to queue  $Q_1$
- At  $Q_1$  job is again served in RR and receives 16 additional milliseconds
  - ▶ If it still does not complete, it is preempted and moved to queue  $Q_2$

